

LAPACK / LAPACKE Eigen

Δρ. Γ. Χαμόδρακας

LAPACK library (1)

- Fortran 90 library
- Provides routines for:
 - solving systems of simultaneous linear equations,
 - least-squares solutions of linear systems of equations,
 - eigenvalue problems,
 - singular value problems,
 - LU, Cholesky, QR, SVD, Schur, generalized Schur matrix factorizations

LAPACK library (2)

- Handles dense and banded but not general sparse matrices
- Supports real and complex matrices in single and double precision
 - i.e. routines with float and double arguments
- <http://www.netlib.org/lapack/>

LAPACKE C interface

- A 2-level C interface to LAPACK Fortran 90 routines
- **High level interface:**
 - handles all workspace memory allocation internally
- Middle level interface:
 - requires the user to provide workspace arrays as in the original Fortran 90 library
- Supports both column-major and row-major matrices

Column-major / Row-major matrices

- 2d-array example

$$\begin{bmatrix} 11 & 12 & 13 \\ 21 & 22 & 23 \end{bmatrix}$$

**Column-major
(Fortran)**

Address	Coords	Value
0	[1,1]	11
1	[2,1]	21
2	[1,2]	12
3	[2,2]	22
4	[1,3]	13
5	[2,3]	23

**Row-major
(C)**

Address	Coords	Value
0	[1,1]	11
1	[1,2]	12
2	[1,3]	13
3	[2,1]	21
4	[2,2]	22
5	[2,3]	23

LAPACKE (cont.d)

- Naming Scheme of the C interface:
 - e.g. dgeev: LAPACK routine for eigenvalue and eigenvector computation of a general matrix
 - LAPACKE_dgeev: high-level interface, LAPACKE_dgeev_work: middle-level interface
- 2d-array arguments:
 - Passed as pointers and not as pointers to pointers
 - Due to this pass an extra int parameter for row-major or col-major storage
 - LAPACK_ROW_MAJOR / LAPACK_COL_MAJOR symbolic constants
 - Row-major layout may require more memory and time

LAPACKE (cont.d)

- Extra argument for leading dimension in matrices:
 - row-major, the number of rows
 - col-major, the number of columns
- Prototypes, symbolic constants, macros, type definitions declared in lapacke.h header file
- lapack_int: integer type defined in lapacke.h

LAPACKE installation

1. Download source from: <http://www.netlib.org/lapack/>
2. `tar xvf lapack-3.6.0.tar`
3. `nano CMakeLists.txt` (within lapack-3.6.0 root dir)
4. Find and edit the following lines (replace OFF with ON):

```
option(LAPACKE "Build LAPACKE" ON)
# LAPACKE has also the interface to some routines from tmglib,
# if LAPACKE_WITH_TMGLIB is selected, we need to add those routines to LAPACKE
option(LAPACKE_WITH_TMGLIB "Build LAPACKE with tmglib routines" ON)
if (LAPACKE_WITH_TMGLIB)
  set(LAPACKE ON)
  if(NOT BUILD_TESTING)
    add_subdirectory(TESTING/MATGEN)
  endif(NOT BUILD_TESTING)
endif(LAPACKE_WITH_TMGLIB)
```

5. `cmake -DBUILD_SHARED_LIBS=ON` (within lapack-3.6.0 root dir)
 6. Run `make` and copy `lib/*` to `/usr/lib`; `include/*` to `/usr/include`
 7. Run `ldconfig`
- Ubuntu, use `apt-get` (version 3.5.0):
 - `sudo apt-get install liblapack-dev checkinstall / sudo apt-get install liblapack-doc`

LAPACKE inverse

- Computation of the inverse of a matrix using the LU factorization
 - First, call `LAPACKE_dgetrf` to fill the `ipiv` pivot indices array.
 - `lapack_int LAPACKE_dgetri (int matrix_layout, lapack_int n, double * a, lapack_int lda, const lapack_int * ipiv)`
- Documentation:
 - http://www.netlib.org/lapack/explore-html/da/d0e/lapacke_dgetri_8c.html
 - http://www.netlib.org/lapack/explore-html/df/da4/dgetri_8f.html

INPUT /OUTPUT PARAMETERS

`matrix_layout`: `LAPACK_ROW_MAJOR` or `LAPACK_COL_MAJOR`

`n`: order of the matrix (num of rows/cols)

`a`: the matrix, size `n*n` (inverse is stored here)

`lda`: the leading dimension of `a == n`

`ipiv`: `int` array, size `== n`

Matrix Multiplication

- LAPACKE incorporates CBLAS (the C interface for Basic Linear Algebra Subprograms)
 - ```
void cblas_dgemm(const enum CBLAS_ORDER Order, const enum CBLAS_TRANSPOSE TransA, const enum CBLAS_TRANSPOSE TransB, const int M, const int N, const int K, const double alpha, const double *A, const int lda, const double *B, const int ldb, const double beta, double *C, const int ldc);
```
- Documentation:
  - [http://www.netlib.org/lapack/explore-html/d7/d2b/dgemm\\_8f.html](http://www.netlib.org/lapack/explore-html/d7/d2b/dgemm_8f.html)
  - <http://www.netlib.org/clapack/CLAPACK-3.1.1.1/BLAS/WRAP/cblas.h>
  - Note the differences for row and col-major arrays (enum and not symbolic constant: CblasRowMajor, CblasColMajor). It is possible to use LAPACK\_COL\_MAJOR, LAPACK\_ROW\_MAJOR

# CBLAS Use

---

1. Download CBLAS code: <http://www.netlib.org/blas/blast-forum/cblas.tgz>
2. `$ tar xvf cblas.tgz`
3. Copy `cblas*.h` to `/usr/include`
4. Create symlink in `/usr/lib` and run `ldconfig`  
`$ sudo ln -sf libblas.so.3 libblas.so`  
`$ sudo ldconfig`
5. Compile with `-lblas` gcc parameter

# LAPACKE eigenproblems

---

- Eigenvalue and eigenvector computation:

- `lapack_int LAPACKE_dgeev( int matrix_layout, char jobvl, char jobvr, lapack_int n, double* a, lapack_int lda, double* wr, double* wi, double* vl, lapack_int ldvl, double* vr, lapack_int ldvr );`

- Documentation:

- [http://www.netlib.org/lapack/explore-html/d9/d28/dgeev\\_8f.html](http://www.netlib.org/lapack/explore-html/d9/d28/dgeev_8f.html)

- [http://www.netlib.org/lapack/explore-html/de/ddd/lapacke\\_dgeev\\_8c.html](http://www.netlib.org/lapack/explore-html/de/ddd/lapacke_dgeev_8c.html)

- Generalised eigenproblem solution:

- `lapack_int LAPACKE_dggeev( int matrix_layout, char jobvl, char jobvr, lapack_int n, double* a, lapack_int lda, double* b, lapack_int ldb, double* alphas, double* alphas, double* betas, double* vl, lapack_int ldvl, double* vr, lapack_int ldvr );`

- Documentation:

- [http://www.netlib.org/lapack/explore-html/d9/d52/dggeev\\_8f.html](http://www.netlib.org/lapack/explore-html/d9/d52/dggeev_8f.html)

- [http://www.netlib.org/lapack/explore-html/de/d27/lapacke\\_dggeev\\_8c.html](http://www.netlib.org/lapack/explore-html/de/d27/lapacke_dggeev_8c.html)

# LAPACKE eigenproblem example (1)

---

```
lapack_int LAPACKE_dgeev(int matrix_layout, char jobvl, char
jobvr, lapack_int n, double* a, lapack_int lda, double* wr,
double* wi, double* vl, lapack_int ldvl, double* vr,
lapack_int ldvr);
```

## INPUT PARAMETERS

matrix\_layout: LAPACK\_ROW\_MAJOR or LAPACK\_COL\_MAJOR

jobvl: **'N'** (do not compute left eigenvectors) / **'V'** compute

jobvr: **'N'** (do not compute right eigenvectors) / **'V'** **compute**

n: order of the matrix (num of rows/cols)

a: the matrix, size n\*n

lda: the leading dimension of a == n

ldvl, ldvr: the leading dimension of left and right  
eigenvector matrices == n

# LAPACK eigenproblem example (2)

---

## OUTPUT PARAMETERS

`wr, wi`: arrays with size at least `n`, containing the real and imaginary parts of the eigenvalues respectively

`vr`: array containing the right eigenvectors, size at least `n*n`  
If `wi[j]== 0` (note that `fabs` function must be used because `wi[j]` is double), corresponding eigenvector is the `j` column of `vr`.

Compilation/Linking example:

```
$gcc -o eigenprob eigenprob.c -llapacke
```

# LAPACKE eigenproblem example (3)

```
#include <stdio.h>
#include <lapacke.h>
#include <math.h>

int main (int argc, const char * argv[])
{
 double a[3][3] = {6,3,-8,0,-2,0,1,0,-3};
 int i,j;
 lapack_int info,n,lda,ldvl,ldvr;
 double *real, *imaginary, *rvecors, *lvecors;
 n = lda = ldvl = ldvr= 3;
 real = LAPACKE_malloc(sizeof(double)*n);
 imaginary = LAPACKE_malloc(sizeof(double)*n);
 rvecors = LAPACKE_malloc(sizeof(double)*n*n);

 info = LAPACKE_dgeev(LAPACK_ROW_MAJOR, 'N', 'V', n, *a, lda, real, imaginary, lvecors, ldvl, rvecors, ldvr);
 for(i=0;i<n;i++)
 {
 printf("Eigenvalue: %lf + %lfi\n",real[i], imaginary[i]);
 printf("EigenvectorTransposed: [");
 if (fabs(imaginary[i])<10e-7) {
 for (j=0;j<n;j++){
 printf("%lf ", rvecors[j*3+i]);
 }
 printf("]\n");
 }
 }
 LAPACKE_free(real);
 LAPACKE_free(imaginary);
 LAPACKE_free(rvecors);
 return(info);
}
```

# LAPACKE solutions (1)

---

$$A = \begin{bmatrix} 6 & 3 & -8 \\ 0 & -2 & 0 \\ 1 & 0 & -3 \end{bmatrix}$$

- Eigenvalues:  
[5, -2, -2] (geometric multiplicity = 1)
- Eigenvectors:

$$\begin{bmatrix} 0,992278 \\ 0 \\ 0,124035 \end{bmatrix}, \begin{bmatrix} 0,707107 \\ 0 \\ 0,707107 \end{bmatrix}, \begin{bmatrix} 0,707107 \\ 0 \\ 0,707107 \end{bmatrix}$$



# LAPACKE solutions (2)

---

$$A = \begin{bmatrix} 0 & 1 & 1 \\ 1 & 0 & 1 \\ 1 & 1 & 0 \end{bmatrix}$$

- Eigenvalues:  
[-1, 2, -1] (geometric multiplicity = 2)
- Eigenvectors:

$$\begin{bmatrix} -0,816497 \\ 0,408248 \\ 0,408248 \end{bmatrix}, \begin{bmatrix} 0,577350 \\ 0,577350 \\ 0,577350 \end{bmatrix}, \begin{bmatrix} 0,226455 \\ -0,792594 \\ 0,566139 \end{bmatrix}$$

# LAPACKE Jacobi SVD

---

- Jacobi Singular Value Decomposition:

- `lapack_int LAPACKE_dgesvj (int matrix_order, char joba, char jobu, char jobv, lapack_int m, lapack_int n, double *a, lapack_int lda, double *sva, lapack_int mv, double *v, lapack_int ldv, double *stat);`

- Documentation:

- [http://www.netlib.org/lapack/explore-html/d1/d5e/dgesvj\\_8f.html](http://www.netlib.org/lapack/explore-html/d1/d5e/dgesvj_8f.html)

- [http://www.netlib.org/lapack/explore-html/d3/d01/lapacke\\_dgesvj\\_8c.html](http://www.netlib.org/lapack/explore-html/d3/d01/lapacke_dgesvj_8c.html)

# LAPACKE Jacobi SVD example (1)

---

## INPUT PARAMETERS

matrix\_layout: LAPACK\_ROW\_MAJOR or LAPACK\_COL\_MAJOR  
joba: **'G'** for a general MxN matrix  
jobu: **'N'** U matrix is not computed  
jobv: **'N'** V matrix is not computed  
m, n: rows, columns of the matrix  
a: the matrix, size m\*n  
lda: the leading dimension of a  
mv: relevant only if V is computed  
stat: double array for internal work, size MAX(6,m+n)

## OUTPUT PARAMETERS

sva: double array, size n for storing singular values  
v: double array for storing matrix V, not referenced in the example

# LAPACKE Jacobi SVD example (2)

---

```
#include <stdio.h>
#include <lapacke.h>
#include <math.h>

int main (int argc, const char * argv[])
{
 double input[3][3] = {1,1,0,0,1,1,1,2,1};
 int i,j;
 lapack_int info,n,ldinput,ldv;
 double *singular, *v, *stat;

 n = ldinput = ldv = 3;
 singular = LAPACKE_malloc(sizeof(double)*n);
 stat = LAPACKE_malloc(sizeof(double)*2*n);

 info = LAPACKE_dgesvj(LAPACK_ROW_MAJOR, 'G', 'N', 'N', n, n, *input, ldinput, singular, 0, v, ldv, stat);
 for(i=0;i<n;i++)
 printf("Singular Value: %.6e\n",singular[i]);

 LAPACKE_free(singular);
 return(info);
}
```

# Eigen Library

---

- Eigen is a C++ template library for linear algebra: matrices, vectors, numerical solvers, and related algorithms.
- <http://eigen.tuxfamily.org/dox/>
- Eigen supports dense and sparse matrices
- Installation:
  - Download and extract the source code: Eigen is a template library (subdirectory Eigen within the extracted directory)
    - Compilation/Linking: `g++ -I /path/to/eigen/ my_program.cpp -o my_program`
  - (Optional) Copy or symlink the header files to `/usr/local/include`
    - Compilation/Linking: `g++ my_program.cpp -o my_program`

# Eigen Library Objects

---

- Eigen provides two kinds of dense objects: mathematical matrices and vectors which are both represented by the template class `Matrix`, and general 1D and 2D arrays represented by the template class `Array`:

```
typedef Matrix<Scalar, RowsAtCompileTime, ColsAtCompileTime,
Options> MyMatrixType;
```

```
typedef Array<Scalar, RowsAtCompileTime, ColsAtCompileTime,
Options> MyArrayType;
```

# Eigen Library Type Parameters

---

- Scalar is the scalar type of the coefficients (e.g., float, double, bool, int, etc.).
- RowsAtCompileTime and ColsAtCompileTime are the number of rows and columns of the matrix as known at compile-time or Dynamic.
- Options can be ColMajor or RowMajor, default is ColMajor.

```
Matrix<double, 6, Dynamic> // Dynamic number of columns (heap allocation)
Matrix<double, Dynamic, 2> // Dynamic number of rows (heap allocation)
Matrix<double, Dynamic, Dynamic, RowMajor> // Fully dynamic, row major (heap)
Matrix<double, 13, 3> // Fully fixed (usually allocated on stack)
```

# Eigen Convenience Typedefs

---

`Matrix<float, Dynamic, Dynamic> <=> MatrixXf`

`Matrix<double, Dynamic, 1> <=> VectorXd`

`Matrix<int, 1, Dynamic> <=> RowVectorXi`

`Matrix<float, 3, 3> <=> Matrix3f`

`Matrix<float, 4, 1> <=> Vector4f`

`Array<float, Dynamic, Dynamic> <=> ArrayXXf`

`Array<double, Dynamic, 1> <=> ArrayXd`

`Array<int, 1, Dynamic> <=> RowArrayXi`

`Array<float, 3, 3> <=> Array33f`

`Array<float, 4, 1> <=> Array4f`



# Eigen Constructors

- **1D**

```
Vector4d v4;
Vector2f v1(x, y);
Array3i v2(x, y, z);
Vector4d v3(x, y, z, w);
VectorXf v5; // empty object
ArrayXf v6(size);

Vector3f v1; v1 << x, y, z;
ArrayXf v2(4); v2 << 1, 2, 3, 4
```

- **2D**

```
Matrix4f m1;
MatrixXf m5; // empty object
MatrixXf m6(nb_rows, nb_columns);

Matrix3f m1; m1 << 1, 2, 3,
 4, 5, 6,
 7, 8, 9;
```

# Eigen Constructors (external arrays)

---

```
float data[] = {1,2,3,4};
Map<Vector3f> v1(data); // uses v1 as a Vector3f object
Map<ArrayXf> v2(data,3); // uses v2 as a ArrayXf object
Map<Array22f> m1(data); // uses m1 as a Array22f object
Map<MatrixXf> m2(data,2,2); // uses m2 as a MatrixXf object
```

rows() function returns the number of rows of a Matrix object  
cols() function returns the number of columns of a Matrix object

size() function returns the total number of coefficients of a Matrix object

# Eigen inverse / determinant

---

- Inverse of a matrix is computed by the inherited member function `inverse()` of the Matrix class template (LU decomposition)
- Determinant of a matrix is computed by the member function `determinant()` of the Matrix class template (LU decomposition)
- Matrix-matrix multiplication is performed with the overloaded `operator*`
- Scalar-matrix multiplication is performed with the `operator *`
- Check the documentation

# Eigenproblem example with Eigen

---

```
#include <iostream>
#include <Eigen/Dense>
#include <Eigen/Eigenvalues>

using std::endl;
using std::cout;

using Eigen::MatrixXd;
using Eigen::EigenSolver;
using Eigen::Map;

int main()
{
 double data[] = {6,0,1,3,-2,0,-8,0,-3};
 Map<MatrixXd> m(data, 3,3);

 EigenSolver<MatrixXd> es(m);

 cout << "The eigenvalues of A are:" << endl << es.eigenvalues() << endl;
 cout << "The matrix of eigenvectors, V, is:" << endl << es.eigenvectors() << endl << endl;

 return 0;
}
```

**Note that matrix data are given with col-major layout**

# Eigen solutions (1)

---

$$A = \begin{bmatrix} 6 & 3 & -8 \\ 0 & -2 & 0 \\ 1 & 0 & -3 \end{bmatrix}$$

- Eigenvalues:

[5, -2, -2] (geometric multiplicity = 1)

- Eigenvectors:

$$\begin{bmatrix} 0,992278 \\ 0 \\ 0,124035 \end{bmatrix}, \begin{bmatrix} -0,707107 \\ 0 \\ -0,707107 \end{bmatrix}, \begin{bmatrix} 0,707107 \\ 0 \\ 0,707107 \end{bmatrix}$$

- Note that  $ev2 = -1*ev3$  (linearly dependent)
- Row 2 of  $ev3$  actually computed as:  $-7.3271e-16$

# Eigen solutions (2)

---

$$A = \begin{bmatrix} 0 & 1 & 1 \\ 1 & 0 & 1 \\ 1 & 1 & 0 \end{bmatrix}$$

- Eigenvalues:  
[-1, 2, -1] (geometric multiplicity = 2)
- Eigenvectors:

$$\begin{bmatrix} -0,816497 \\ 0,408248 \\ 0,408248 \end{bmatrix}, \begin{bmatrix} 0,577350 \\ 0,577350 \\ 0,577350 \end{bmatrix}, \begin{bmatrix} 0 \\ -0,707107 \\ 0,707107 \end{bmatrix}$$

- Note that LAPACKE\_ev3  $\cong$  -0.277350\*ev1 + 0.960769\*Eigen\_ev3
- Linear combination of independent eigenvectors

# Produced Eigenvectors

---

- When the geometric multiplicity of an eigenvalue is greater than 1, eigenvectors produced by LAPACK and Eigen may be different and not linearly dependent.
- In that case, produced eigenvectors are a linear combination of  $N$  independent eigenvectors, where  $N$  is the geometric multiplicity of the eigenvalue.

# Jacobi SVD example with Eigen

```
#include <iostream>
#include <Eigen/Dense>
#include <Eigen/Core>
#include <Eigen/SVD>

using std::endl;
using std::cout;
using Eigen::MatrixXd;
using Eigen::JacobiSVD;
using Eigen::Map;

int main()
{
 double data[] = {1,0,1,1,1,2,0,1,1};
 Map<MatrixXd> m(data, 3,3);

 JacobiSVD<MatrixXd> svd(m, Eigen::ComputeFullU);
 JacobiSVD<MatrixXd>::SingularValuesType singular = svd.singularValues();

 cout << "The singular values of A are:" << svd.singularValues() << endl;
 for (int i = 0; i < singular.rows(); i++)
 cout << "Singular Value" << i << ":" << singular(i) << endl;

 return 0;
}
```



# Eigen Functors

---

- `eigenvalues()`, `eigenvectors()` and `singularValues()` functions in the previous examples return objects of the `Matrix` template class (typedef'd)
- The `Matrix` template class is a functor. Specifically, the matrix coefficient accessors and mutators are provided through the overloaded parenthesis operator:

```
MatrixXd m(2,2);
m(0,0) = 3;
m(1,0) = 2.5;
m(0,1) = -1;
m(1,1) = m(1,0) + m(0,1);
```

# Generalized Eigenproblem

---

- Eigen does not fully support the solution of a Generalized Eigenproblem
- Eigenvectors are not computed
- In order to resolve this, LAPACKE may be interfaced with Eigen
- LAPACKE must be installed

# Generalized Eigenproblem example

```
#include <iostream>
#include <Eigen/Dense>
#include <lapacke.h>

using std::endl;
using std::cout;
using Eigen::MatrixXd;

bool GEP(MatrixXd& A, MatrixXd& B, MatrixXd& v, MatrixXd& lambda);

int main()
{
 MatrixXd A = MatrixXd::Random(4,4);
 MatrixXd B = MatrixXd::Random(4,4);
 MatrixXd V(4,4); // Contains N=4 Eigenvectors (4 rows each)

 /* Contains N=4 Eigenvalues. each eigen value has a real and an imaginary part
 (columns 0 and 1) and a denominator (column 2) */
 MatrixXd lambda(4,3);

 GEP(A,B,V,lambda);
 cout << lambda << endl;
 cout << V << endl;
}

bool GEP(MatrixXd& A, MatrixXd& B, MatrixXd& v, MatrixXd& lambda)
{
 int N = A.cols(); // Number of columns of A and B. Number of rows of v.
 if (B.cols() != N || A.rows() != N || B.rows() != N)
 return false;

 v.resize(N,N);
 lambda.resize(N, 3);

 int LDA = A.outerStride(); int LDB = B.outerStride(); int LDV = v.outerStride(); int INFO = 0;

 double * alphas = lambda.col(0).data();
 double * alphas_i = lambda.col(1).data();
 double * betas = lambda.col(2).data();

 INFO = LAPACKE_dggeev(LAPACK_COL_MAJOR, 'N', 'V', N, A.data(), LDA, B.data(), LDB, alphas, alphas_i, betas, 0, LDV, v.data(), LDV);

 return INFO==0;
}
```