

Collaborative Recommendation Techniques

Ioannis Chamodrakas, Ph.D.
Prof. Ioannis Emiris

Problem Statement

$U = \{u_1, \dots, u_i, \dots, u_N\}$ a set of N users

$I = \{i_1, \dots, i_j, \dots, i_M\}$ a set of M items

$R: U \times I \rightarrow R_o$, a rating function

Usually $R_o \subseteq N = \{0, 1, 2, \dots\}$

Problem: Predict unknown $R(u_i, i_j)$ when values
 $R(u_k, i_l)$ known for many $k \in \{1, \dots, N\}$, $l \in \{1, \dots, M\}$.

Recommend best or few highest rated items
(usually up to 5) for user u_i .

Memory based approach

Derive ratings from similar users (or items) according to past ratings, e.g. some NNs according to some similarity function.

In practice, neighborhood size $P \in \{20-50\}$.

$$R(u,i) = z * \sum \text{sim}(u,v) * R(v,i)$$

or, rather:

$$R(u,i) = R(u) + z * \sum \text{sim}(u,v) * (R(v,i) - R(v))$$

$v \in \text{Neighborhood}(u)$, z = normalizing factor

(e.g. $1 / \sum |\text{sim}(u,v)|$), $R(u)$ =average rating of u .

Memory based: similarity

- A good similarity function is Cosine distance:

$$\text{sim}(u,v) = u \cdot v / (\|u\| \cdot \|v\|)$$

- Every user is represented by a vector in M-dim space of items, with item ratings as coordinates.

Non-rated items usually treated as zero.

- Issue: lack of rating more similar to disliking than liking item.
- Other approaches: Rounding, or assign to unrated items a constant, adjusted experimentally to minimize error.
- Our approach: Normalization (next slide).

Memory based: unrated

Two ways to treat unrated items

- **Normalization (our approach):**

$$R'(u, i) = R(u, i) - R(u),$$

$R(u)$ = average rating of u .

- **Rounding (cutoff):**

E.g. $R_o = \{1, 2, 3, 4, 5\}$: $R'(u, i) = 0$ if $R(u, i) \leq 2$, $R'(u, i) = 1$ if $R(u, i) > 2$.

Model based approach

- **models users based on their past ratings, then uses models to predict unknown ratings.**
 - **Cluster Models**
 - **Bayesian Models**
 - **Latent semantic models**

Recommendation Algorithms

NN (by Cosine LSH):

1. Find P neighbors with Cosine LSH (slides 1.nnLSH) binary-repeated range search, $20 \leq P \leq 50$.
2. Calculate $R(u,i)$ from ratings of P closest users.
3. Recommend 5 top-rated items.

Binary-repeated search: multiply radius by $\frac{1}{2}$ or 2, then bisect relevant interval.

k-Clustering:

1. Cluster the users: optimize k through Silhouette.
2. Calculate $R(u,i)$ from ratings of users in same cluster.
3. Recommend 5 top-rated items.

Algorithm Validation

- **Idea: partition dataset into training (known input) and validation subsets; predict ratings in validation subset and compare with true ratings.**

- **Mean Absolute Error:**

$$\text{MAE} = (1/J) * \sum |R_j - P_j|, \quad 1 \leq j \leq J$$

J=#(item,user) pairs in validation subset i.e. ratings known and predicted: R_j =actual rating, P_j =predicted rating

***f*-fold cross-Validation**

- ***f*-fold validation of recommendation algorithms:**
 1. **Split the rating dataset into f equal-size subsets randomly. E.g. uniformly pick J/f pairs, $f-1$ times; or pairs uniformly mapped in $\{1...f\}$.**
 2. **Repeat f times (folds): each subset used exactly once for validation, the others used as training data to predict ratings; then calculate Mean Absolute Error (MAE).**
 3. **Return average MAE (or other combination) from the folds.**

[Wikipedia]